

AI PROMPT ENGINEERING LIBRARY

A Comprehensive Collection of Production-Ready Prompts for Career, Strategy, Analytics, and Automation

TOOL 1: JOB SEARCH & CAREER ACCELERATION FRAMEWORK

Prompt 1.1: Resume Diagnostic Before Rewriting

Act as a senior career strategist, ATS-aware recruiter and honest editor.
I am attaching my resume. Do not rewrite it yet. First analyse it and return:

1. My strongest skills and evidence for each.
2. Weak or unclear sections.
3. Missing information you need from me.
4. Whether my career narrative is clear in one sentence.
5. The roles I appear strongest for today.
6. A table with: issue, why it matters, priority score (0-10), and the exact information needed to fix it.

Rules:

- Do not infer achievements that are not written.
- Distinguish missing evidence from missing skill.
- Ask up to five questions after the analysis, then wait.

Prompt 1.2: Find Roles That Fit the Resume

Based on my resume and the career direction below, identify 10 roles I should target.
Target direction: [AI generalist / AI orchestrator / function-specific AI-first role]
Location or remote requirement: []
Minimum compensation: []
Industries I prefer or want to avoid: []
Experience constraints: []

For each role, return:

- likely job titles and title variations,
- why I fit,
- the biggest gap,
- keywords to search,
- a fit score from 0-100,
- whether I should apply now, apply after a small portfolio project, or skip.

Do not recommend a role only because the title contains "AI". Match the actual work.

Prompt 1.3: Tailor the Resume to One Job Description

Act as an ATS resume strategist and hiring manager for the attached job description.
Inputs:

- My original resume: attached
- Target job description: attached or pasted below

Tasks:

1. Extract the core responsibilities, must-have skills, preferred skills and repeated keywords from the JD.
2. Compare them with my resume and create a gap table with a priority score from 0-10.
3. Rewrite my resume for this specific role while preserving only truthful facts.
4. Improve the summary, experience bullets, skills order and project framing.

5. Quantify impact only when the source resume contains enough evidence.
6. Return an ATS-friendly version with simple headings and no tables, icons or graphics.
7. Add a change log explaining every major edit.
8. Flag any claim that needs my confirmation before use.

Output order:

- A. Fit summary
- B. Gap analysis table
- C. Tailored resume
- D. Keyword coverage checklist
- E. Claims requiring confirmation

Prompt 1.4: 30-Day Gap-Closing Plan

Using the gap analysis between my resume and this job description, create a realistic 30-day learning and proof-of-work plan.

For each gap, include:

- whether it is a knowledge gap, evidence gap or experience gap,
- priority from 0-10,
- the smallest useful learning task,
- one portfolio artifact I can create,
- how to describe the artifact honestly on my resume or LinkedIn,
- time estimate,
- success check.

Limit the plan to [X] hours per week. Prioritise visible proof of skill over passive course completion.

Prompt 1.5: Strict Mock Interviewer

Act as a strict but fair interviewer for the attached job description. Use my attached resume to personalise the interview.

Interview rules:

- Ask one question at a time and wait for my spoken answer.
- Ask 10-15 questions in total.
- Mix behavioural, role-specific, problem-solving and resume-deep-dive questions.
- Challenge vague claims and ask follow-ups when evidence is missing.
- Do not coach me during the interview and do not reveal ideal answers.
- Keep the difficulty at [intermediate / senior / leadership] level.

After the final question, give a full debrief with:

1. Overall score out of 10.
2. Scores for structure, clarity, evidence, role knowledge, confidence and concision.
3. My three strongest answers.
4. My three weakest answers and better answer structures.
5. Missed opportunities from my resume.
6. A 3-day practice plan.

Prompt 1.6: Investor Grilling Variant

Act as a sceptical seed-stage investor. I am attaching my pitch deck.

Ask one hard question at a time about market, customer pain, differentiation, traction, distribution, business model, team and risks.

Do not accept broad answers; ask for evidence, assumptions and numbers.
After 12 questions, provide an investment-readiness score, the five biggest unanswered risks and the exact slides or data I should improve.

Prompt 1.7: Tough Client Meeting Variant

Act as the most demanding prospective client for the attached proposal.
Ask one question at a time about scope, timelines, ownership, measurement, dependencies, cost, risks and what happens when the plan fails.
After 10 questions, create a client-objection sheet with: objection, what they are really worried about, evidence needed, and a concise response.

TOOL 2: REAL-TIME RESEARCH & IDEA GENERATION (X / GROK)

Prompt 2.1: Engagement-Led Research on X

Act as a real-time social research analyst.
Search X for public conversations about [topic] during [date range].
Prioritise posts by meaningful engagement – replies, reposts, likes and views – rather than recency alone.

Analyse:

1. The top recurring problems or questions.
2. The language people use to describe the problem.
3. Common complaints about existing solutions.
4. Desired outcomes and willingness-to-pay signals.
5. High-engagement creators and communities discussing the topic.
6. Contrarian or emerging viewpoints.

Return:

- a table of themes with supporting links,
- short compliant excerpts rather than long copied text,
- engagement indicators,
- a confidence level for each pattern,
- 20 content or product ideas derived from the evidence.

Separate observed evidence from your interpretation. Do not invent posts, handles or metrics.

Prompt 2.2: Turn Raw Research into Ranked Ideas

I am pasting research findings below.
Generate 20 ideas for [content / product / campaign / workshop].
Score each idea from 0-10 on:

- audience pain intensity,
- evidence in the research,
- differentiation,
- ease of execution,
- business value.

Then rank the top five. For each top idea, state the exact evidence that supports it, the riskiest assumption and the smallest test I can run within seven days.

Research:

[PASTE RESEARCH]

TOOL 3: MULTI-MODEL AI COUNCIL & DECISION ENGINE

Prompt 3.1: Council Question — Shared Prompt for All Models

You are one member of an independent expert council evaluating the decision below.
Decision: [decision to make]
Context: [background, constraints, objective]
Evidence: [paste or attach source material]

Analyse independently. Do not try to agree with other models.
Return:

1. Your recommendation.
2. The three strongest reasons.
3. The three biggest risks or failure modes.
4. Assumptions that could change the answer.
5. Missing evidence.
6. A low-cost test or next step.
7. Confidence from 0-100 and why.

Prompt 3.2: Convene the Judge — Synthesis Prompt

Act as the chairperson judging multiple AI responses to the same decision.
Do not choose the most confident or longest answer. Evaluate the reasoning and evidence.

Produce:

- A. Points of consensus.
- B. Important disagreements and what causes them.
- C. Unique insights that appear in only one response.
- D. Unsupported claims or hallucination risks.
- E. The best combined recommendation.
- F. Conditions under which the recommendation should change.
- G. A decision table with option, upside, downside, evidence, risk and confidence.
- H. The next action a human should take before committing.

Model responses:
[PASTE ALL RESPONSES]

Prompt 3.3: Second Opinion / Red-Team Pass

Assume the current recommendation is wrong.
Red-team it from the perspective of a sceptical operator, customer, finance leader and domain expert.
Identify hidden assumptions, second-order effects, implementation risks and evidence that would falsify the recommendation.
End with a revised recommendation or explain why the original still holds.

TOOL 4: PRESENTATION & PITCH DECK GENERATION (KIMI)

Prompt 4.1: Create Story and 10-Slide Outline

Act as a presentation strategist.
Using the research below, design a 10-slide narrative for [audience] that helps them [decision or action].

For each slide, provide:

- slide title written as a takeaway,
- the one message the audience must remember,
- supporting evidence or data,
- recommended visual (chart, comparison, timeline, process, image or table),
- speaker note,
- source reference.

Narrative requirements:

1. Start with the audience's problem or decision.
2. Build a clear argument, not a list of facts.
3. Use one main idea per slide.
4. End with an explicit recommendation and next step.
5. Flag weak evidence instead of filling gaps.

Research:

[PASTE RESEARCH]

Prompt 4.2: Generate Deck in Kimi Slides

Create a fully editable 10-slide presentation using the outline below.

Audience: []

Objective: []

Tone: [premium / analytical / energetic / minimal]

Visual direction: [describe colours, typography, image style, spacing and references]

Brand requirements: [logo, colours, fonts, do-not-use list]

Design rules:

- 16:9 widescreen.
- One clear takeaway per slide.
- Large readable text and generous whitespace.
- Use editable charts, diagrams and tables where possible.
- Use high-quality relevant imagery; avoid generic AI robots and visual clutter.
- Keep all claims traceable to the sources in the outline.
- End with a concise action slide.

Outline:

[PASTE OUTLINE]

Prompt 4.3: Replicate a Reference Style Responsibly

I am uploading [2-5] reference slide images.

First analyse the design system: colour palette, font characteristics, title hierarchy, spacing, grid, image treatment, shapes, chart style and recurring motifs.

Then apply the same design logic to my new content without copying logos, proprietary illustrations or brand-owned text.

Keep the deck visually consistent and fully editable.

TOOL 5: EXCEL & GOOGLE SHEETS AUTOMATION (NUMEROUS.AI)

Prompt 5.1: Generate Spreadsheet Formula from Plain English

I want the values in cells F2, G2, H2, I2 and J2 combined into one line with a single space between each value.

Give me the exact Google Sheets formula and explain it in one sentence.

Handle blank cells without creating repeated spaces.

Prompt 5.2: Design a Reliable Bulk-Classification Prompt

Classify the text into exactly one of these labels: [LABEL 1], [LABEL 2], [LABEL 3], [OTHER].

Definitions:

- [LABEL 1]: [definition + example]
- [LABEL 2]: [definition + example]
- [LABEL 3]: [definition + example]

Rules:

- Return only the label, with no explanation.
- If the text is ambiguous or does not match, return OTHER.
- Do not create new labels.

TOOL 6: DATA ANALYSIS & EXECUTIVE DASHBOARDS (CLAUDE)

Prompt 6.1: Executive Data Analysis & Interactive Dashboard

Act as a senior data scientist and strategy consultant.

I am attaching a CSV / spreadsheet.

Work in this order:

1. Audit the data: columns, data types, missing values, duplicates, outliers and limitations.
2. Identify the 4-5 strategic questions a C-suite leader would ask.
3. Select one governing question that should organise the analysis.
4. Analyse the data and explain the method in plain English.
5. Build an interactive dashboard that answers the governing question and the supporting questions.
6. Use [brand colours / visual style] while keeping accessibility and readability high.
7. Add filters only when they help a real decision.
8. Create an executive summary with key insights, business implications and recommended actions.
9. Include a limitations section and do not imply causation from correlation.

Quality rules:

- Never fabricate missing values.
- Show the denominator and units for every metric.
- Verify calculations and reconcile totals.
- Flag any result that is sensitive to data quality or assumptions.

Prompt 6.2: Ask for a Governing Question First

Before analysing the dataset, propose five possible governing questions.

For each, explain:

- why an executive would care,
- which fields can answer it,
- which fields are missing,
- what decision it could inform.

Recommend the strongest question and wait for my approval before building charts.

TOOL 7: ACCELERATED LEARNING & SYNTHESIS (NOTEBOOKLM)

Prompt 7.1: Build a Source-Grounded Learning Notebook

Topic: [topic]
Learning goal: [what I need to understand or do]
Current level: [beginner / intermediate / advanced]
Time available: []

Use only the sources in this notebook.

Create:

1. A concept map of the topic.
2. The 20% of concepts that unlock 80% of understanding.
3. A step-by-step learning path.
4. Common misconceptions.
5. Practical examples connected to my role: [role].
6. Five questions that test real understanding, not memorisation.
7. A list of source gaps where the notebook cannot answer confidently.

Cite the supporting source for every important claim.

Prompt 7.2: Custom Instructions for an Audio Overview

Create an audio overview for a busy [profession / learner type].
Use [English / Hindi / bilingual] language.
Focus on practical intuition and examples, not technical jargon.

Structure it as:

- why this matters,
- three core ideas,
- one realistic example,
- common mistake,
- five-minute action I can take today.

Keep it approximately [X] minutes.

Prompt 7.3: Interactive Question During Audio

Pause and explain [concept] as if I am a non-technical [role].
Use one analogy, one work example and one counterexample.
Then ask me one question to check whether I understood it.

Use only the notebook sources and tell me when the sources are insufficient.

Prompt 7.4: Generate a Retrieval Practice Pack

From the notebook sources, create:

- 15 flashcards,
- 10 scenario-based multiple-choice questions,
- 5 short-answer questions,
- an answer key with source references,
- a one-page revision sheet.

Make the distractors plausible. Test application and judgement, not only definitions.

TOOL 8: END-TO-END PRODUCT, AUTOMATION & AI TEAMMATES

Prompt 8.1: No-Code MVP Product Brief (Lovable, Replit, AI Studio)

Act as a product manager and no-code solution architect.
Problem: [describe a repeated real-world problem]
User: [specific user]
Current workaround: [how they solve it today]
Desired outcome: []
Constraints: [time, budget, data, platform, technical level]

Create:

1. A one-sentence problem statement.
2. The job-to-be-done and top user pains.
3. Assumptions that must be validated.
4. A smallest viable workflow – not a feature-heavy app.
5. A PRD with user stories and acceptance criteria.
6. Recommended no-code / AI tool stack with reasons.
7. Data model and integrations.
8. Privacy, permission and failure considerations.
9. A build sequence for a first working prototype in [X] days.
10. A user test script and success metrics.

Avoid features that do not directly solve the core problem.

Prompt 8.2: Automation Blueprint (n8n, Zapier, Make)

Act as an automation architect.
Turn the process below into a reliable automation blueprint.

Process today: [describe the manual steps]
Start condition / trigger: []
Inputs and systems: []
Desired output: []
Volume and frequency: []
Exceptions: []

Return a table with:

- step,
- trigger or input,
- action,
- tool/system,
- data passed,
- validation check,
- failure path,
- human approval required.

Then recommend the smallest version to automate first.
Include duplicate prevention, logging, retries, privacy and a manual fallback.
Do not automate a step whose decision criteria are not yet clear.

Prompt 8.3: AI Employee Specification (Assistants, Agents & Internal Systems)

Design an AI teammate for this role: [role / job to be done].

Define:

1. Mission and measurable outcomes.
2. Tasks it may perform independently.
3. Tasks that require human approval.
4. Knowledge sources and update frequency.
5. Tools and permissions it needs.

6. Standard operating procedure for each core task.
7. Input and output formats.
8. Quality checks and evaluation rubric.
9. Safety, privacy and compliance guardrails.
10. Escalation rules and failure recovery.
11. Memory: what it should remember and what it must not store.
12. A 14-day pilot plan with success metrics.

End with a one-page "AI employee card" containing role, responsibilities, tools, guardrails and escalation contact.